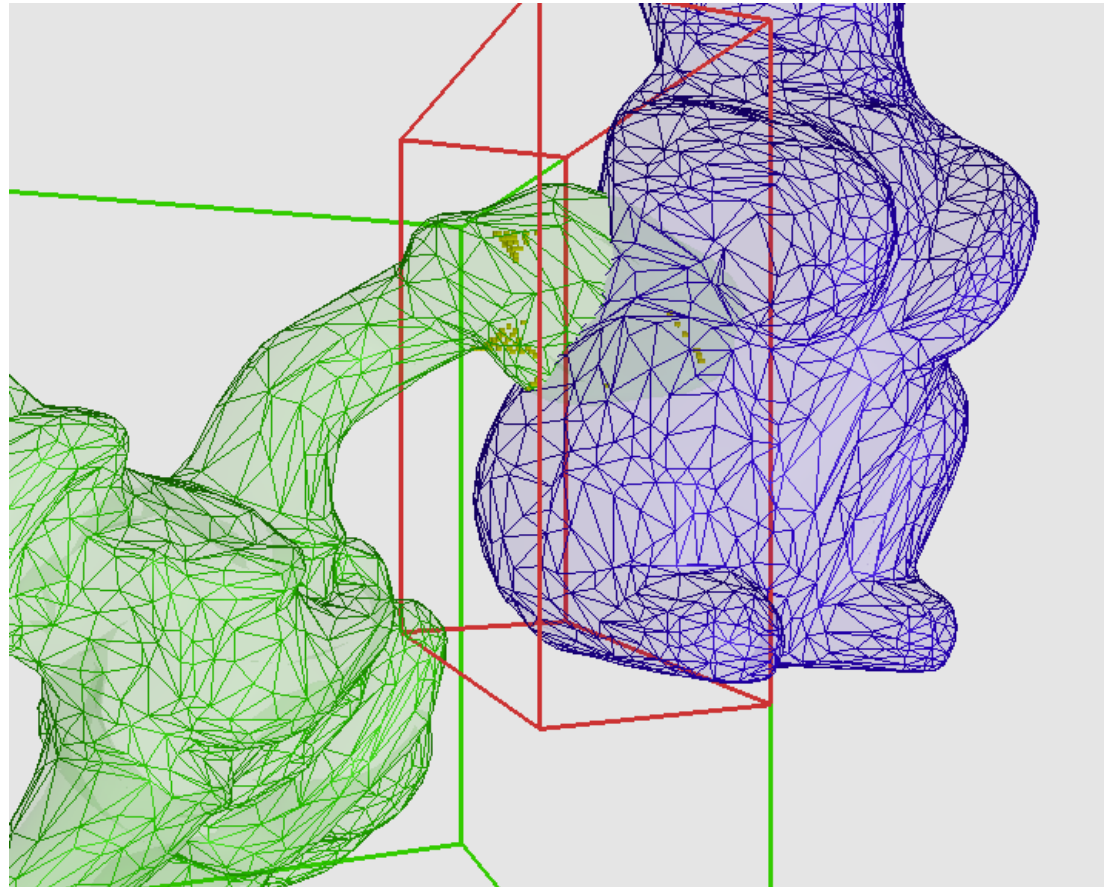




# *Image-Space Collision Detection*





# *Acknowledgements*

- Parts of this slide set are courtesy of Bruno Heidelberger, ETH Zurich.



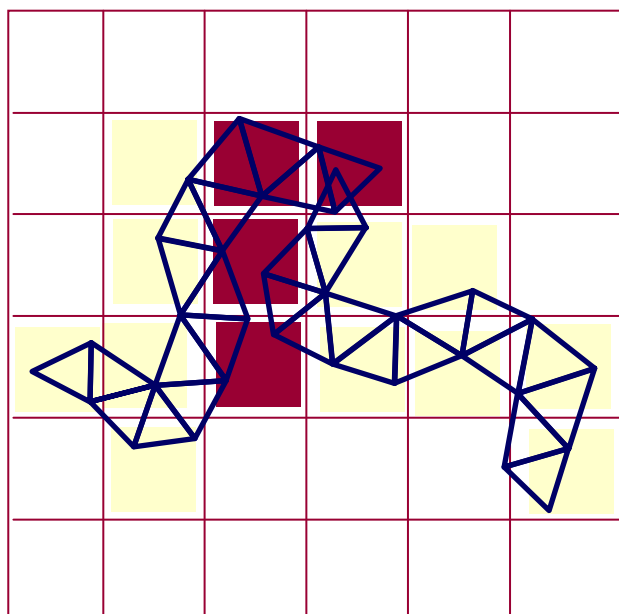
# *Outline*

- motivation
- algorithms
- performance
- application
- discussion



# *Graphics Hardware for 2D Collision Detection*

- rendering corresponds to placing all object primitives into the according cell (pixel) of a uniform rectangular 2D grid (frame buffer)
- rendering determines all pixels in a frame buffer affected by the object
- at each pixel position, information can be processed (color, depth, stencil)

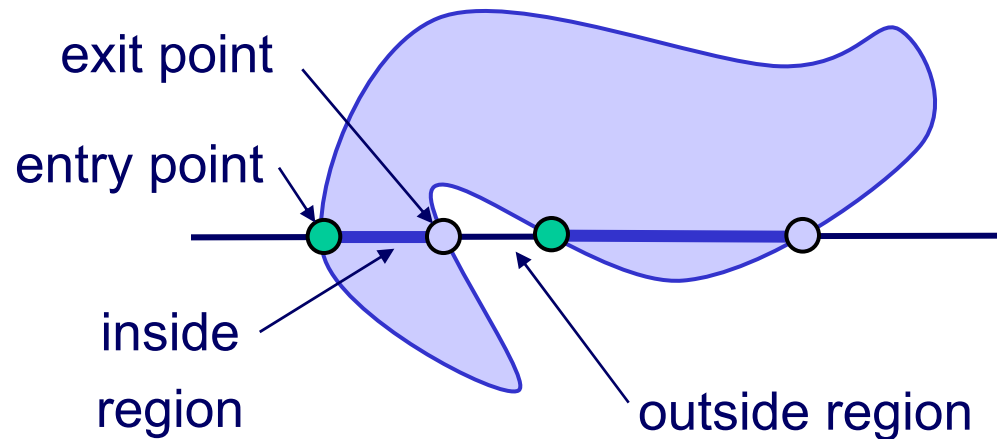


 stencil value 1     stencil value 2

- Kenneth Hoff, UNC
- stencil-buffer for collision detection
- clear stencil buffer
- increment stencil buffer for each rendered object
- intersection for stencil buffer value larger 1



# *Closed Objects*

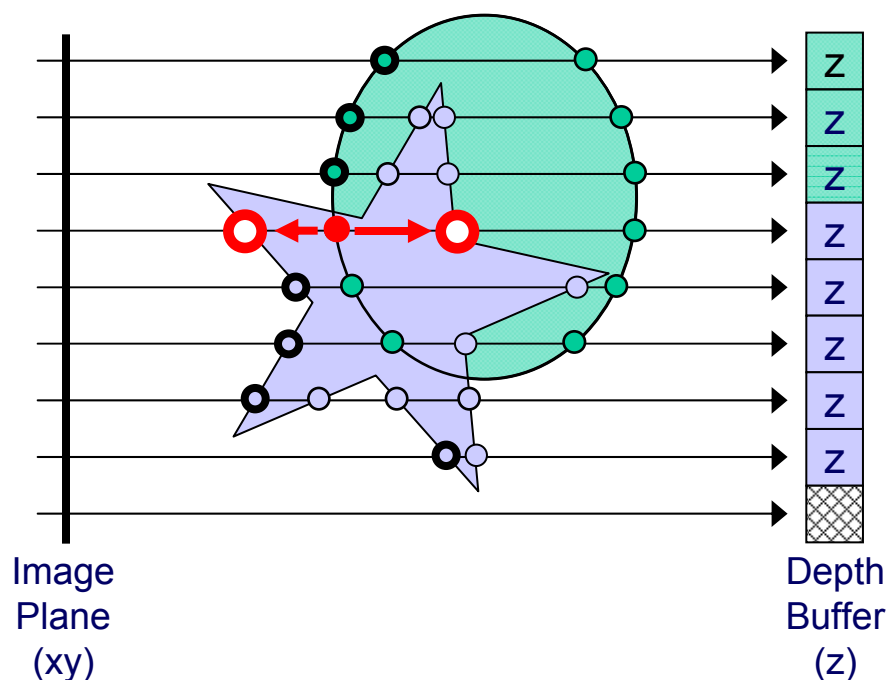


- number of entry points equals the number of exit points
- in case of convex objects, one entry point and one exit point
- inside and outside are separated by entry or exit point
- entry point is at a front face
- exit point is at a back face
- front and back faces alternate



# *Collision Detection with Graphics Hardware*

- exploit rasterization of object primitives for intersection test
- benefit from graphics hardware acceleration





# *Collision Detection with Graphics Hardware*

## **Idea**

- computation of entry and exit points can be accelerated with graphics hardware
- computation corresponds to rasterization of surface primitives
- all object representations that can be rendered are handled
- parallel processing on CPU and GPU

## **Challenges**

- restricted data structures and functionality

## **Drawbacks**

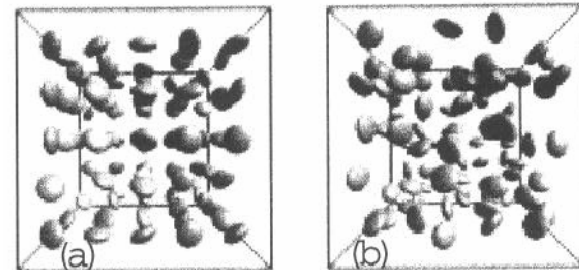
- approximate computation of entry and exit points



# Early approaches

**[Shinya, Forgue 1991]**

image-space collision detection for  
*convex objects*

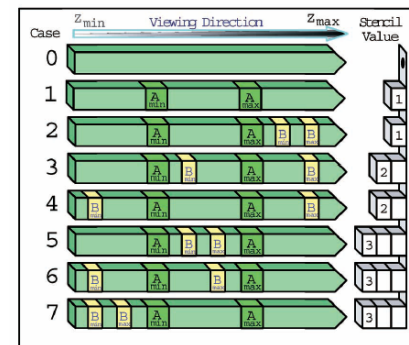


**[Myszkowski, Okunев, Kunii 1995]**

collision detection for *concave objects*  
with limited depth complexity

**[Baciu, Wong 1997]**

hardware-assisted collision detection for  
*convex objects*







# *More approaches*

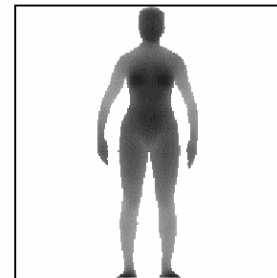
**[Lombardo, Cani, Neyret 1999]**

intersection of *tool* with *deformable tissue*  
by rendering the interior of the tool



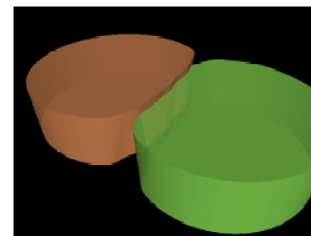
**[Vassilev, Spanlang, Chrysanthou 2001]**

image-space collision detection applied to  
cloth simulation and *convex avatars*



**[Hoff, Zaferakis, Lin, Manocha 2001]**

proximity tests and penetration  
depth computation, *2D*

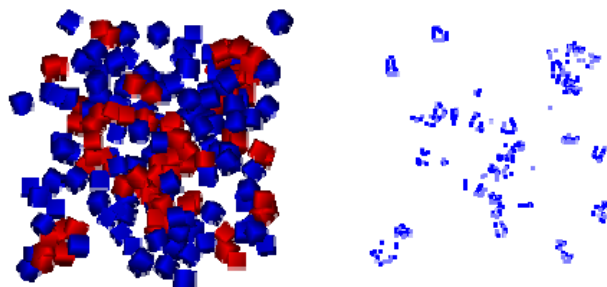




# *Recent approaches*

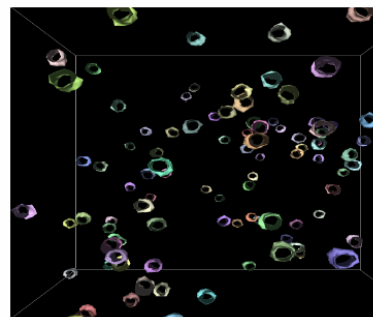
**[Knott, Pai 2003]**

intersection of edges with surfaces



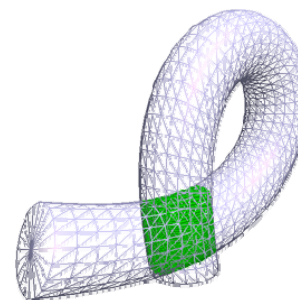
**[Govindaraju, Redon, Lin, Manocha 2003]**

object and sub-object pruning based on occlusion queries



**[Heidelberg, Teschner 2004]**

explicit intersection volume and self-collision detection based on LDIs

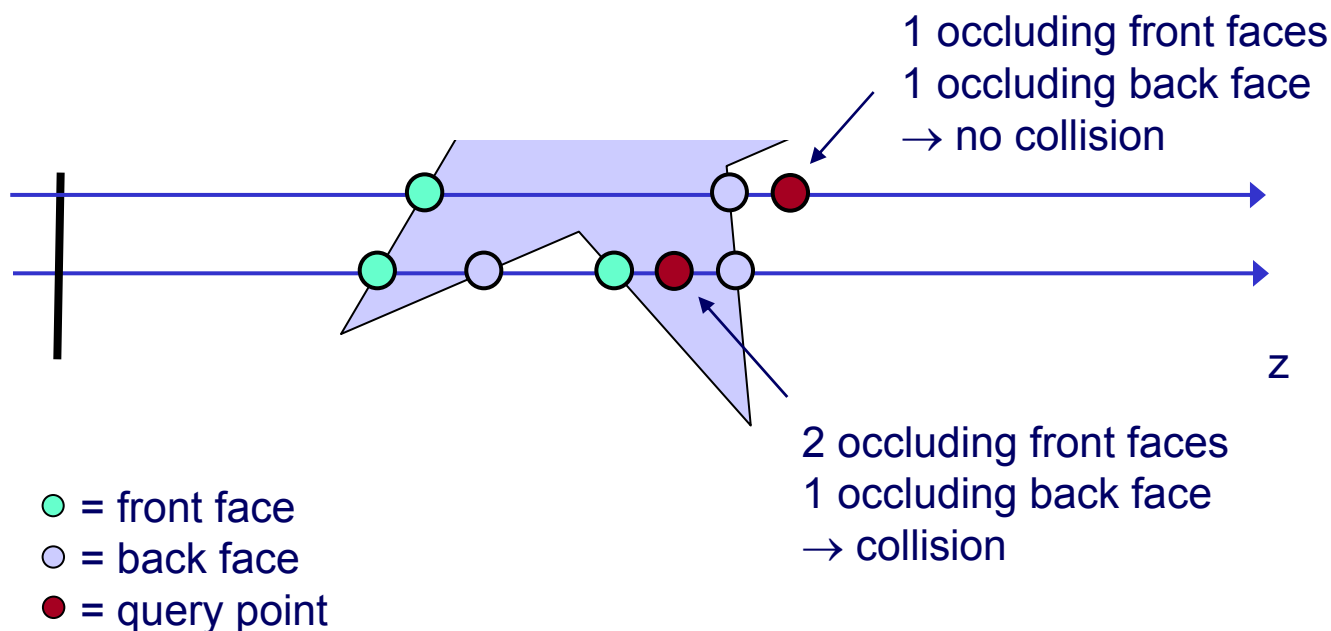




# Image-Space Collision Detection

[Knott, Pai 2003]

- render all query objects (e. g. edges) to depth buffer
- count the number  $f$  of front faces that occlude the query object
- count the number  $b$  of back faces that occlude the query object
- iff  $f - b == 0$  then there is no collision





# *Image-Space Collision Detection*

- clear depth buffer, clear stencil buffer
- render query objects to depth buffer
- disable depth update
- render front faces with stencil increment
  - if front face is closer than query object, then stencil buffer is incremented
  - depth buffer is not updated
  - result: stencil buffer represents number of occluding front faces
- render back faces with stencil decrement
  - if back face is closer than query object, then stencil buffer is decremented
  - depth buffer is not updated
  - result: stencil buffer represents difference of occluding front and back faces
- stencil buffer not equal to zero → collision



# *Image-Space Collision Detection*

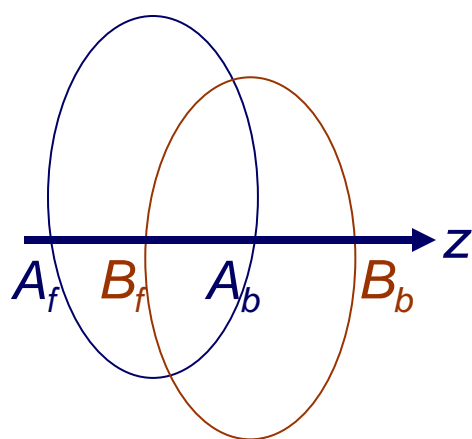
- works for objects with closed surface
- works for n-body environments
- works for query objects that do not overlap in image space
- numerical problems if query object is part of an object
  - offset in z-direction required
- [Video]



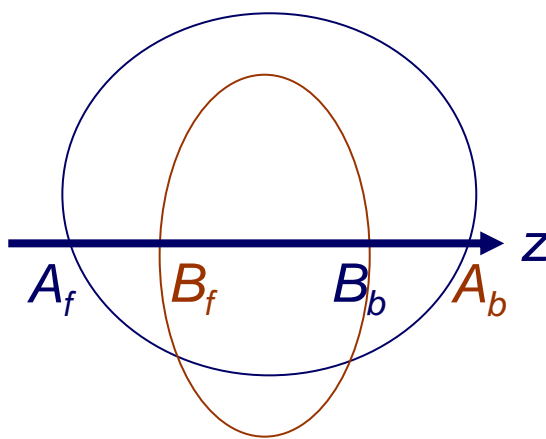
# Image-Space Collision Detection

[Baciu 2000]

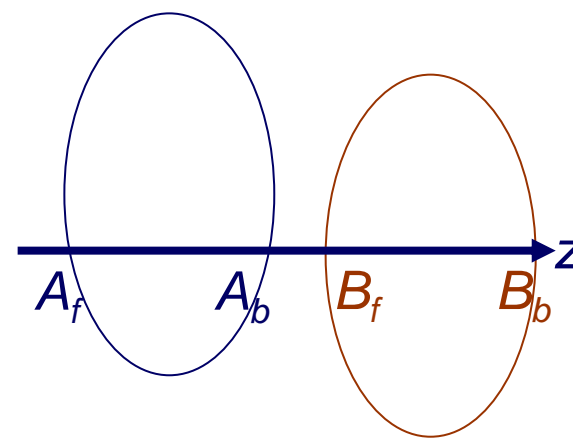
- RECODE – REndered COLLision DETection
- works with pairs of closed convex objects A and B
- one or two rendering passes for A and B
- algorithm estimates overlapping z intervals per pixel



collision



collision



no collision



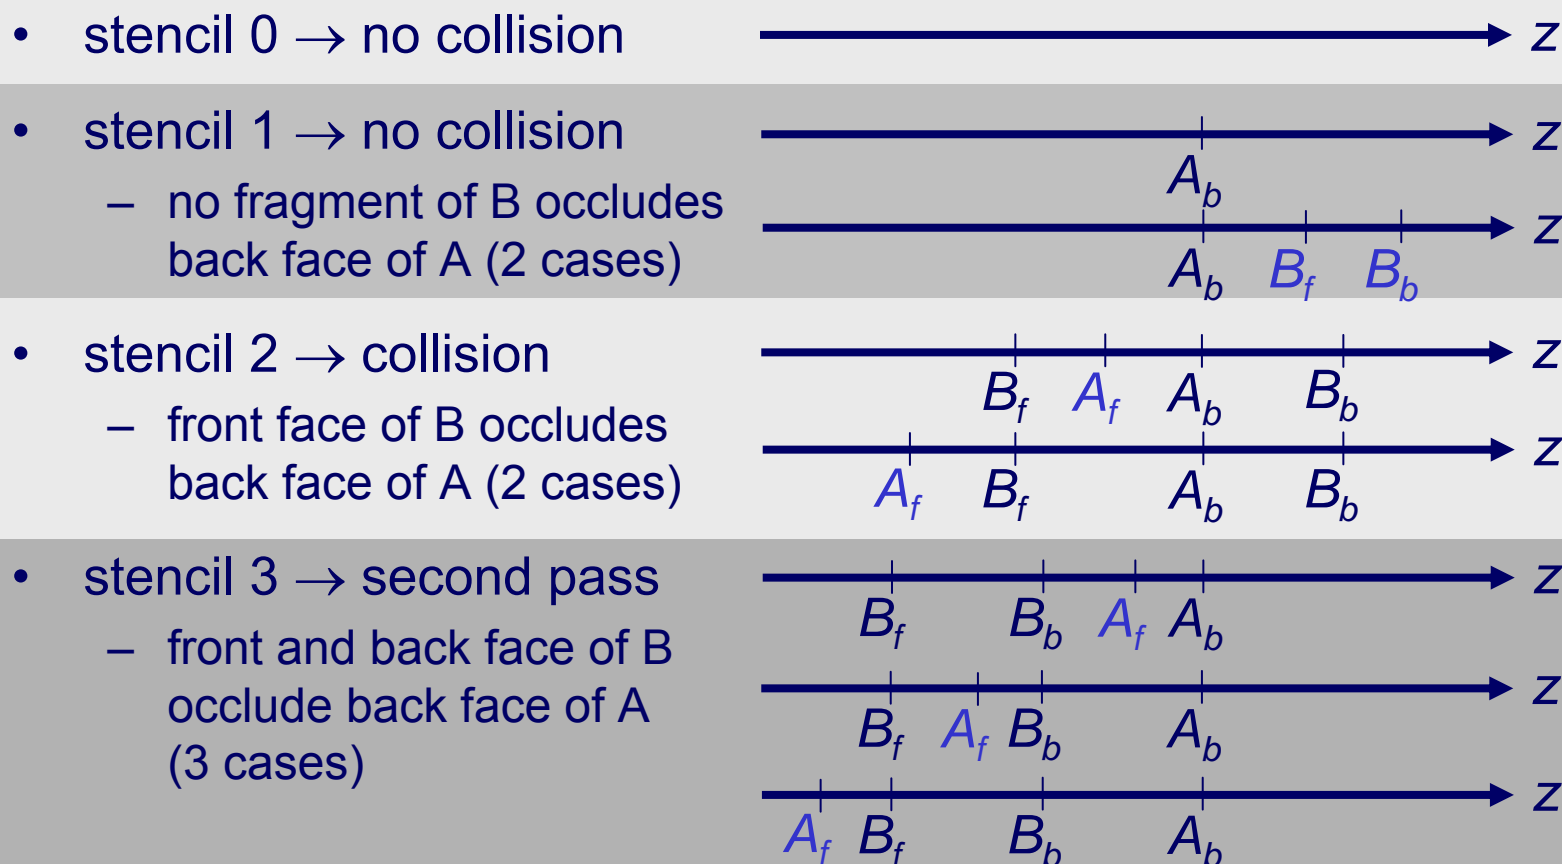
# *First Rendering Pass*

- clear depth buffer
- clear stencil buffer
- enable depth update
- render back faces of A with stencil increment
  - if nothing has been rendered  $\rightarrow$  stencil=0
  - if something has been rendered  $\rightarrow$  stencil=1
  - depth buffer contains depth of back faces of A
- disable depth update
- render B with stencil increment
  - if stencil==1 and B occludes back face of A  $\rightarrow$  stencil+=1
  - depth buffer is not updated
  - stencil-1 = number of faces of B that occlude A



# First Rendering Pass

- first pass collision query







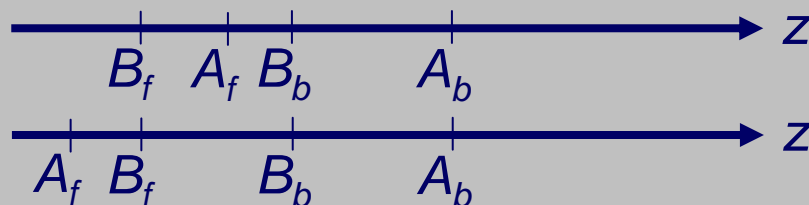
# Second Rendering Pass

- render back faces of object B, count occluding faces of A
  - corresponds to first pass with A and B permuted
  - only 3 cases based on the result of the first rendering pass

- stencil 1 → no collision
  - no fragment of A occludes back face of B (1 case)



- stencil 2 → collision
  - front face of A occludes back face of B (2 cases)



- done



# Second Rendering Pass

[Myszkowski 1995]

- render front faces of object A, count occluding faces of B
  - corresponds to first pass, front faces are rendered instead of back faces
  - only 3 cases based on the result of the first rendering pass

- stencil 3 → no collision
  - front and back face of B occlude front face of A



- stencil 2 → collision
  - front face of B occludes front face of A



- stencil 1 → collision
  - no fragment of B occludes front face of A



- done



# *Image-Space Collision Detection for Concave Objects [Myszkowski 1995]*

- collision detection for pairs of concave objects  
A and B with limited depth complexity (number of entry/exit points)
- faces have to be sorted with respect  
to the direction of the orthogonal projection (e. g. BSP tree)
- objects are rendered in front-to-back or back-to-front order
- alpha blending is employed:  
$$\text{color}_{\text{framebuffer}} = \text{color}_{\text{object}} + \alpha \cdot \text{color}_{\text{framebuffer}}$$
- color of A is zero, color of B is  $2^{k-1}$ ,  
 $k$  is the number of bits in the frame buffer,  
 $\alpha = 0.5$



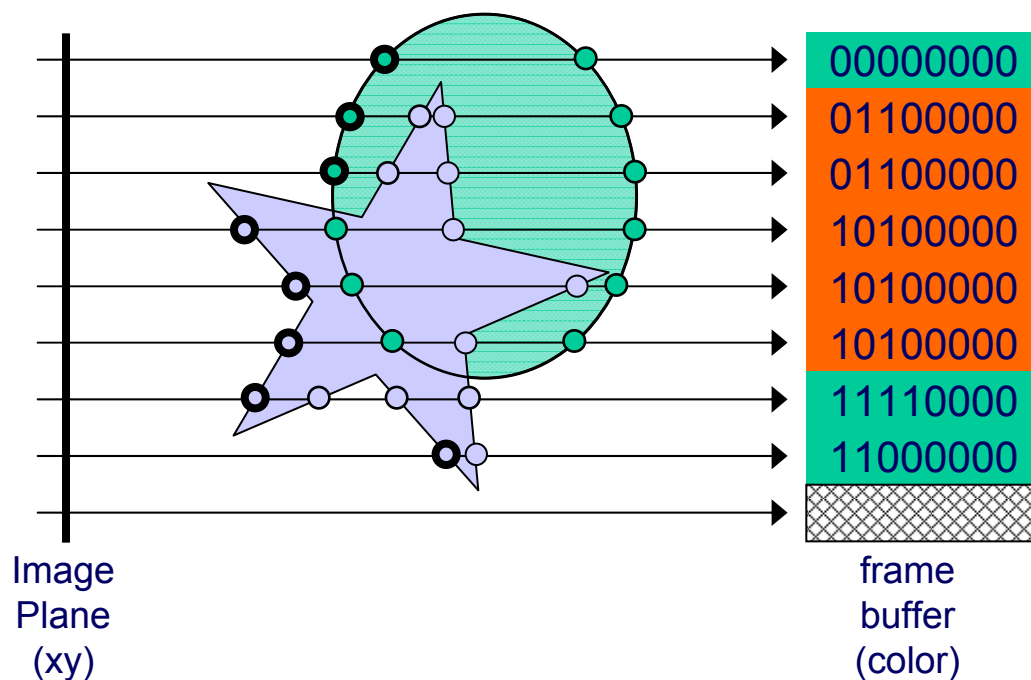
# *Image-Space Collision Detection for Concave Objects*

- example:  $k = 8$
- color A = 0, color B =  $2^7$
- sequence of faces  $B_1 A_1 A_2 B_2 B_3 B_4$  rendered back to front:
  - $c_{fb} = 00000000_2$
  - render  $B_4$ :  $c_{fb} = 2^7 + \alpha \cdot c_{fb} = 10000000_2 + 0.5 \cdot 00000000_2 = 10000000_2$
  - render  $B_3$ :  $c_{fb} = 10000000_2 + 0.5 \cdot 10000000_2 = 11000000_2$
  - render  $B_2$ :  $c_{fb} = 10000000_2 + 0.5 \cdot 11000000_2 = 11100000_2$
  - render  $A_2$ :  $c_{fb} = 00000000_2 + 0.5 \cdot 11100000_2 = 01110000_2$
  - render  $A_1$ :  $c_{fb} = 00000000_2 + 0.5 \cdot 01110000_2 = 00111000_2$
  - render  $B_1$ :  $c_{fb} = 10000000_2 + 0.5 \cdot 00111000_2 = 10011100_2$
- resulting bit sequence represents order of faces of A (0) and B (1)
- odd number of adjacent zeros or ones indicates collision



# Image-Space Collision Detection for Concave Objects

- example:





# *Image-Space Collision Detection*

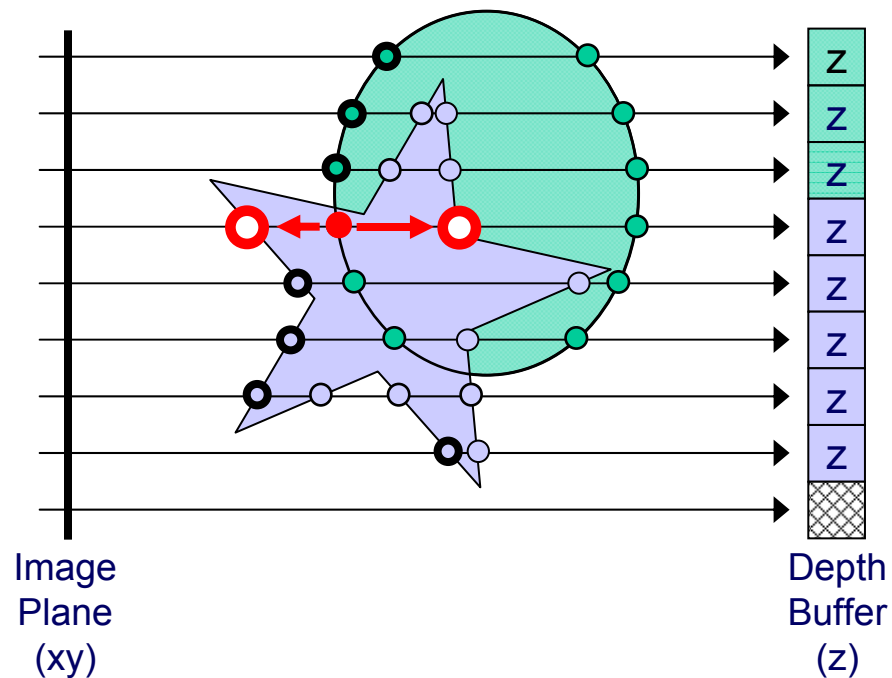
## *[Heidelberg 2003]*

- works with pairs of closed arbitrarily-shaped objects
- three implementations
  - $n+1$  hardware-accelerated rendering passes  
where  $n$  is the depth complexity of an object
  - $n$  hardware-accelerated rendering passes
  - 1 software rendering pass
- three collision queries
  - intersection volume (based on intersecting  $z$  intervals)
  - vertex-in-volume test
  - self-collision test
- basic idea and implementation for convex objects  
has been proposed by Shinya / Forgue in 1991



# *Collision Detection with Graphics Hardware*

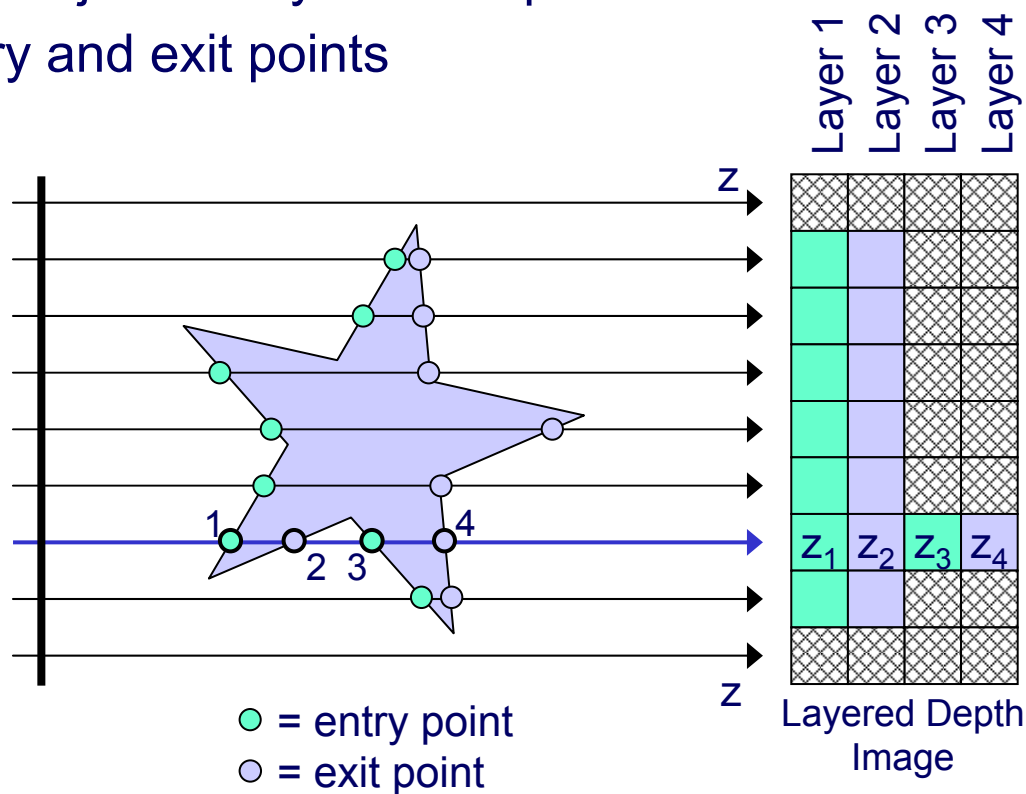
- exploit rasterization of object primitives for intersection test
- benefit from graphics hardware acceleration





# *Layered Depth Image*

- compact, volumetric object representation [Shade et al. 1998]
- represents object as layers of depth values
- stores entry and exit points



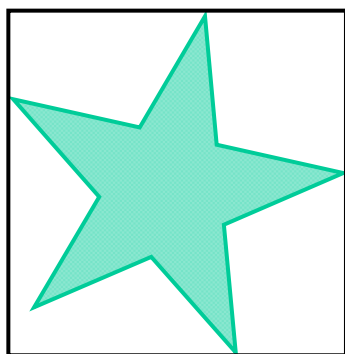




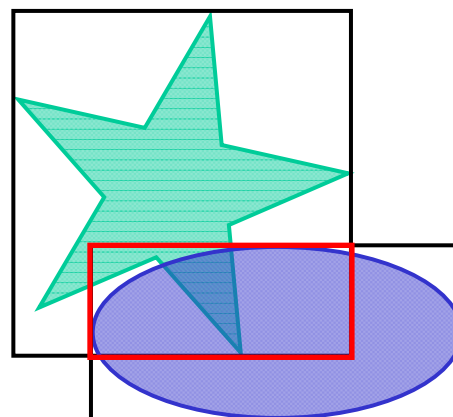
# *Algorithm Overview*

Algorithm consists of 3 stages:

Stage 1: Check for bounding box intersection



a) Very fast detection of trivial “no collision” cases

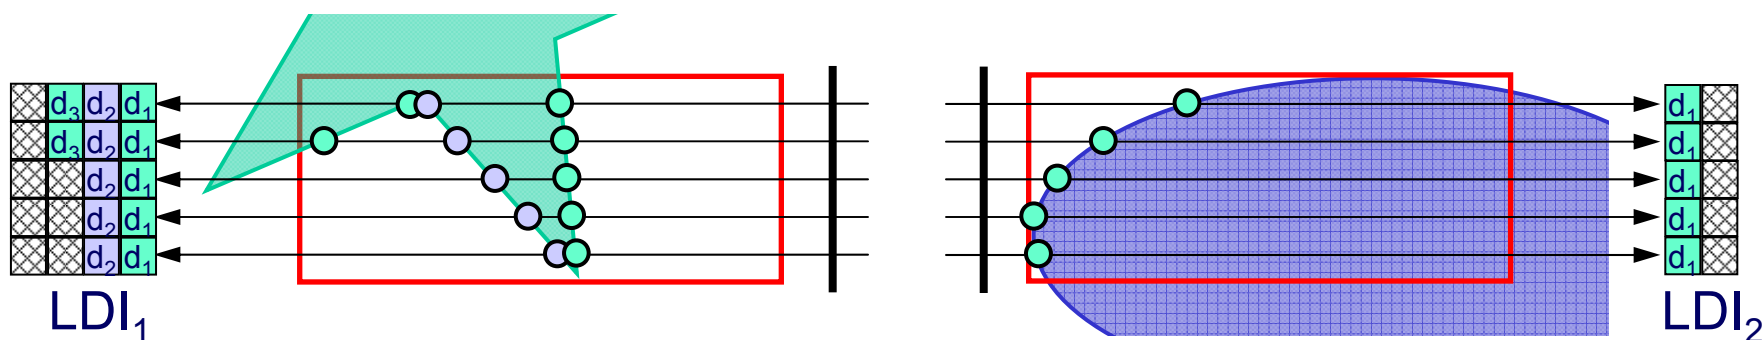


b) Overlapping area defines volume of interest (Vol) for step 2 & 3



# Algorithm Overview

Stage 2: Generate the layered depth images (LDI)



Step 3: Perform the collision tests

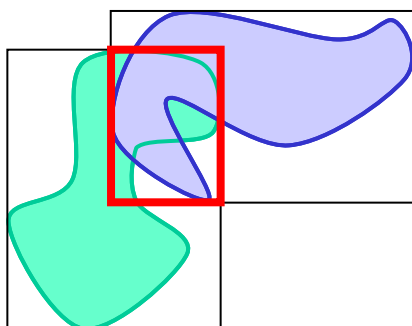
- a) test object primitives of one object against LDI of the other
- b) combine both LDI to get overlapping volume
- c) self-intersection test



# Algorithm Overview

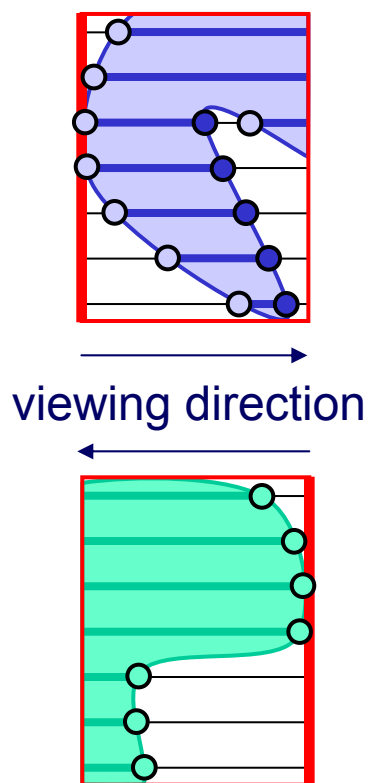
## Stage 1

Volume-of-interest



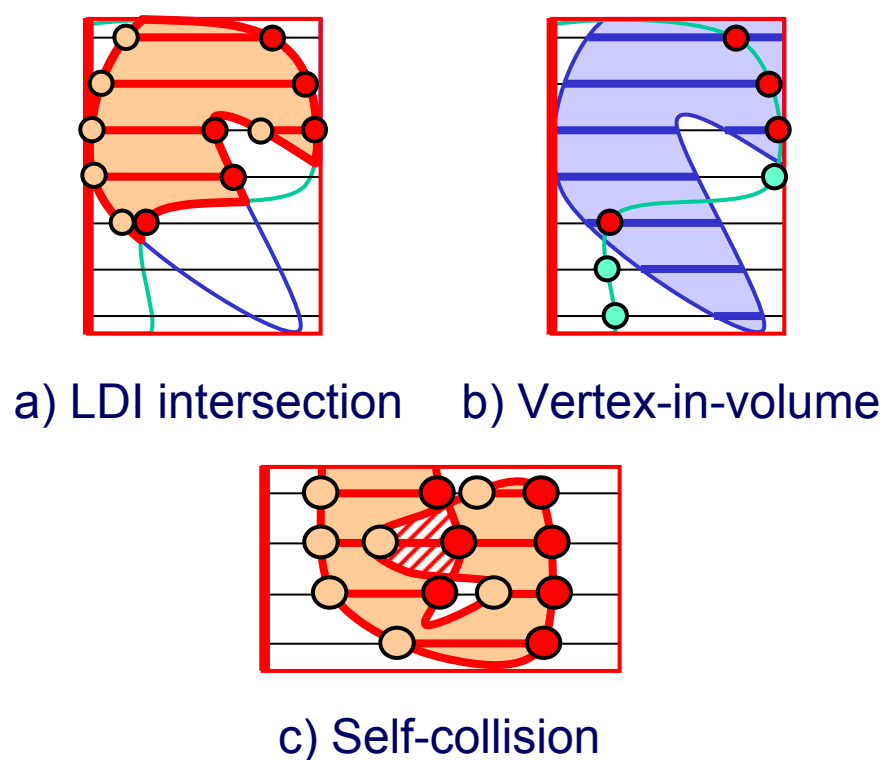
## Stage 2

LDI generation



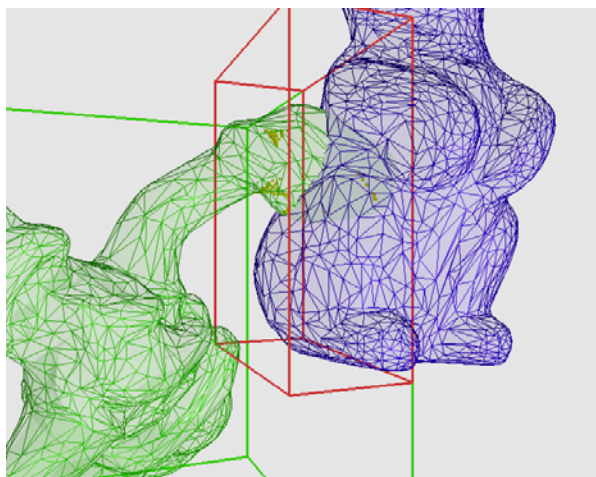
## Stage 3

Collision query

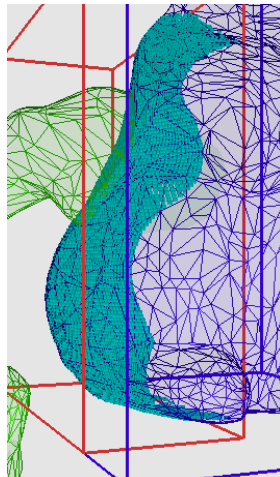




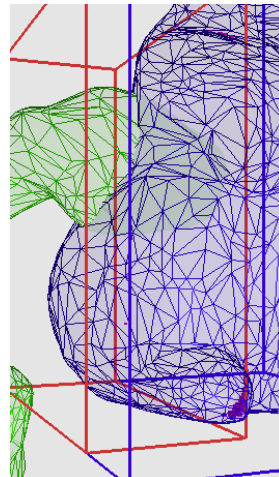
# Algorithm Overview



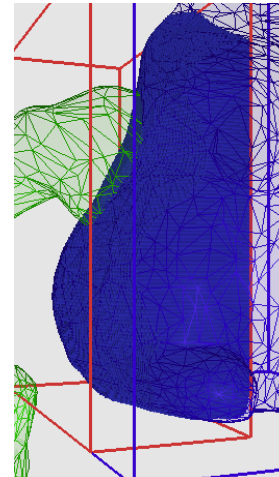
volume of interest Vol



layer 1



layer 2



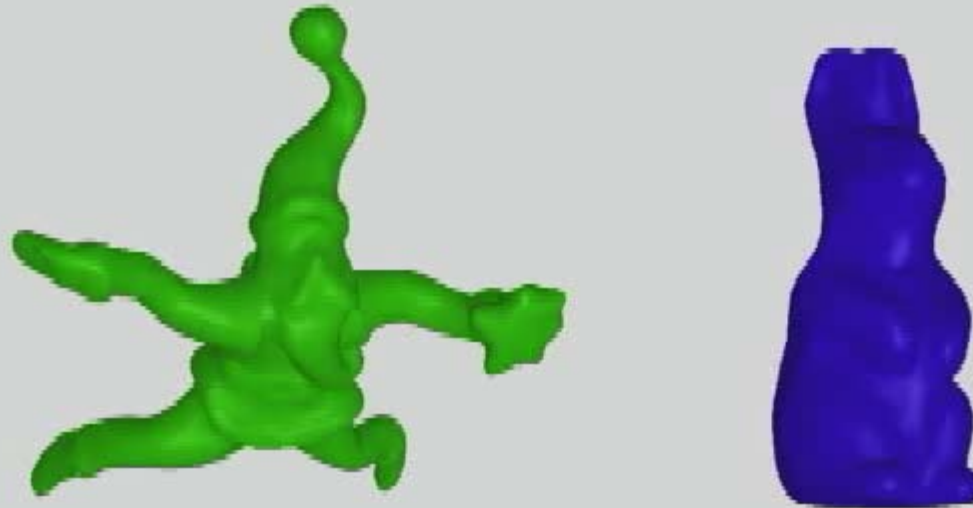
volume



collision  
queries



# *Algorithm Overview*



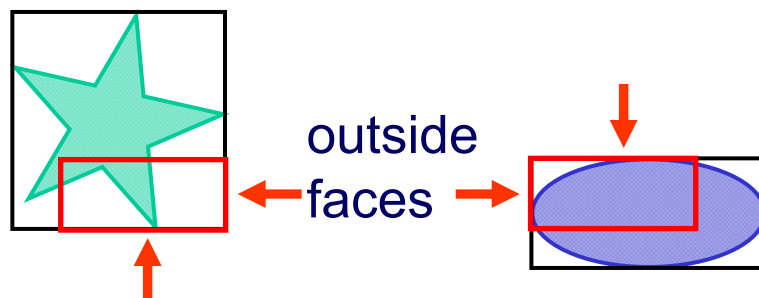
**Real-Time Volumetric Intersections  
of Deforming Objects**



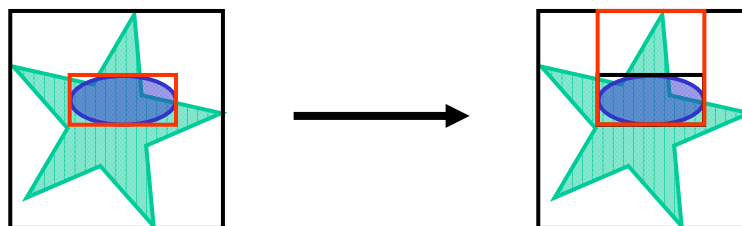
# Volume of Interest

$$\text{Vol} = \text{BoundingBox}(\text{Object 1}) \cap \text{BoundingBox}(\text{Object 2})$$

1. evaluation of trivial rejection test:  $\text{Vol} == \emptyset \rightarrow$  no collision!
2. choice of *opposite* render directions for LDI generation



possible enlargement of Vol to guarantee valid directions



**outside faces are outside the object**

**-> guarantees that first intersection point is an entry point**

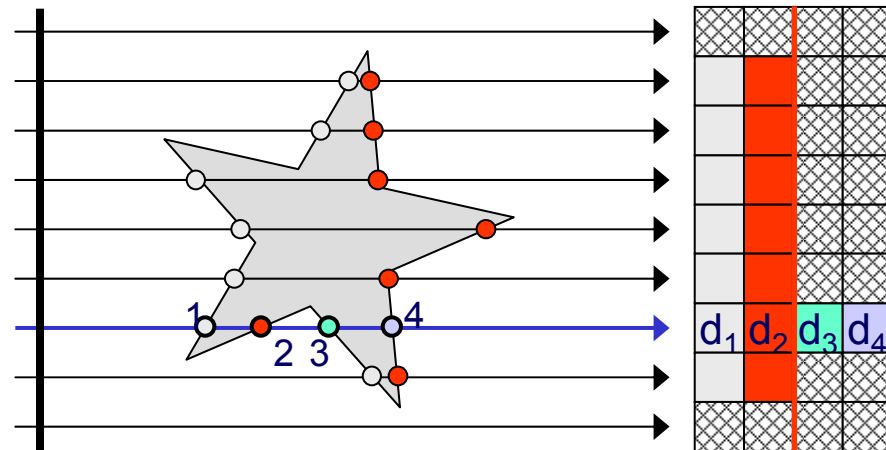


# *LDI Generation on the GPU*

## *Depth Peeling*

- object is rendered once for each layer in the LDI
- two separate depth tests per fragment are necessary:
  - fragment must be **farther** than the one in the previous layer ( $d_2$ )
  - fragment must be **the nearest** of all remaining fragments ( $d_3$  &  $d_4$ )

example: *pass #3*



→ second depth test is realized using shadow mapping  
extended depth-peeling approach [Everitt 2001]



# Shadow Mapping

## Idea:

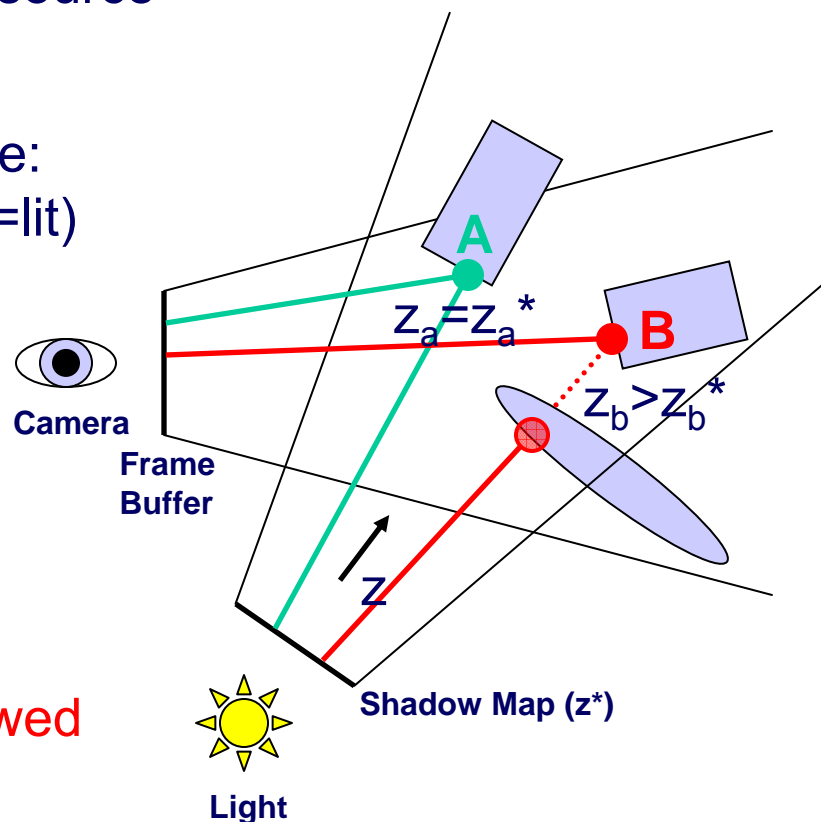
- for each fragment to be rendered:  
check if it is visible from the light source

## Algorithm:

- render scene from the light source:  
store all distances to the visible (=lit) fragments in a “shadow map”
- render scene from the camera:  
compare the distance  $z$  of each fragment to the light with the value  $z^*$  in the shadow map:

$z = z^* \rightarrow$  fragment is lit

$z > z^* \rightarrow$  fragment is shadowed







# *Shadow Mapping as Depth Test*

## Differences to regular depth test:

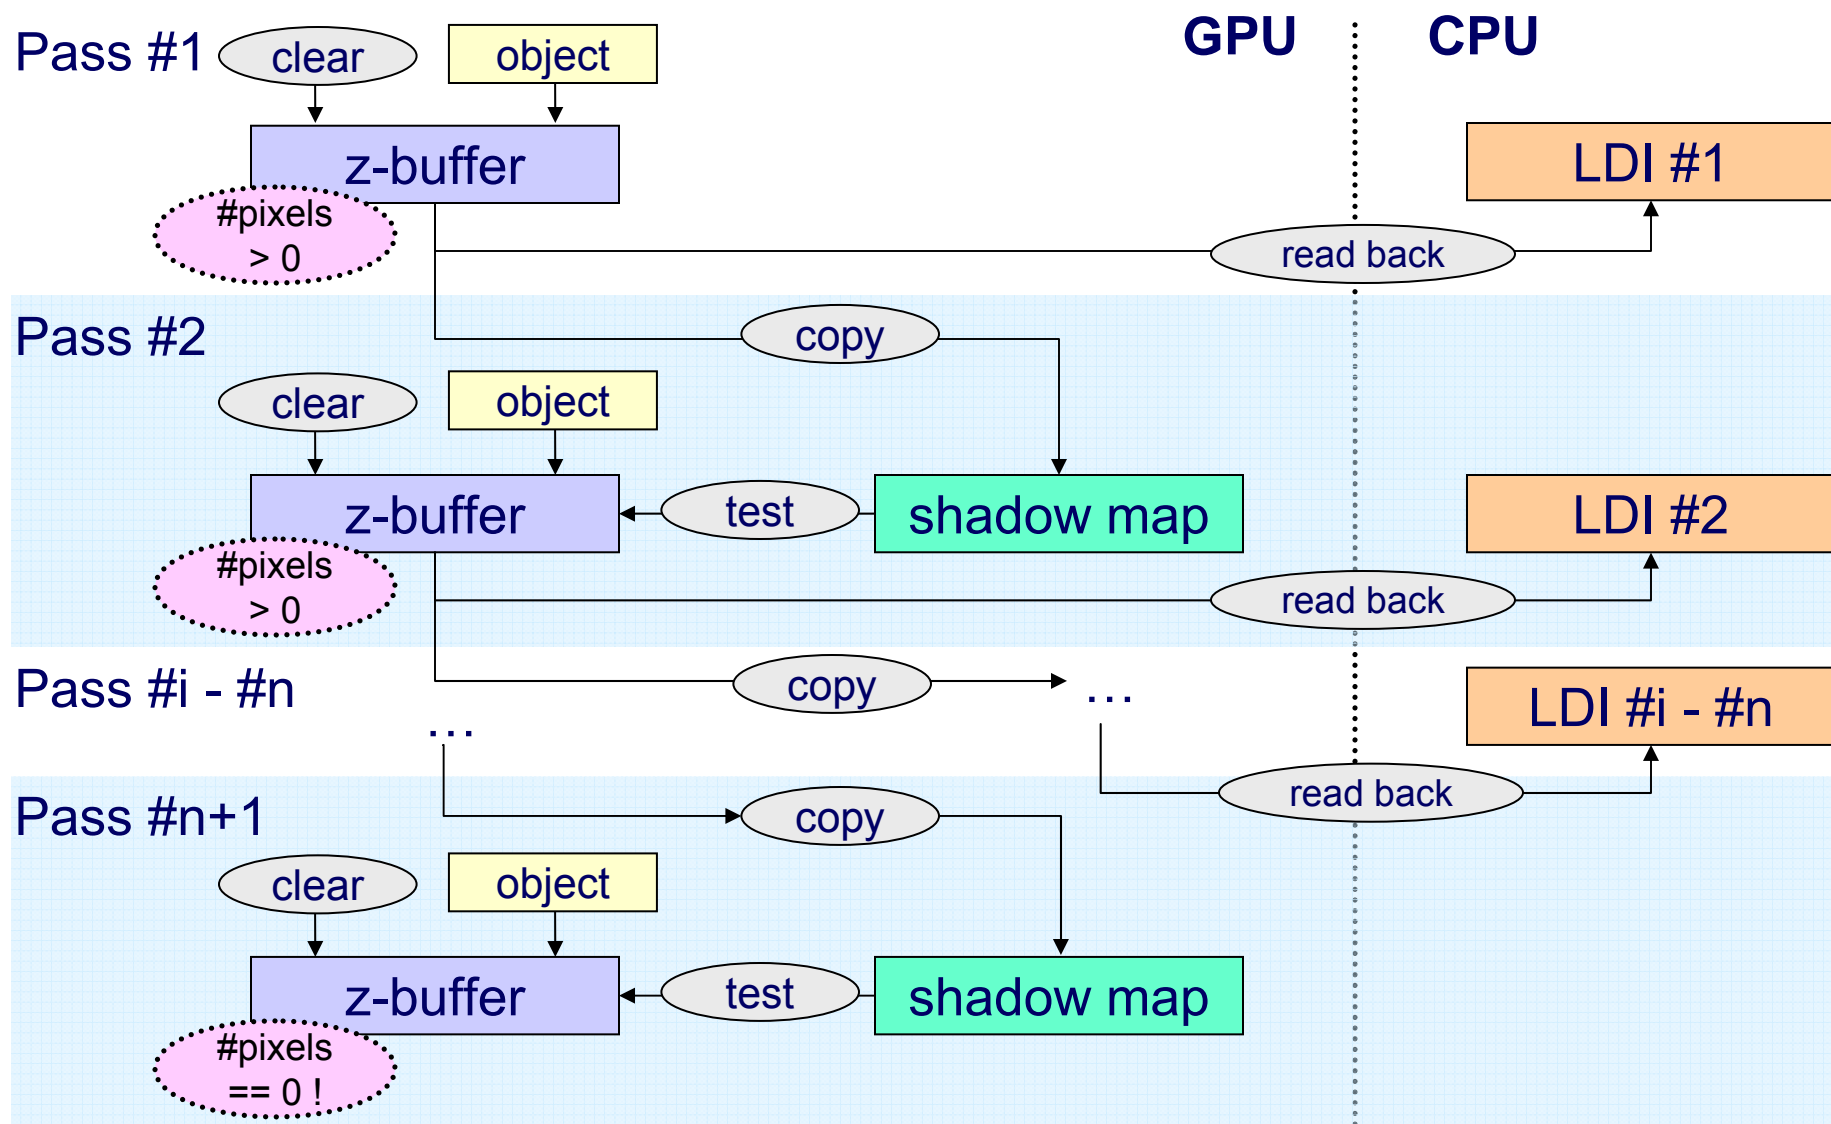
- shadow mapping depth test is not tied to camera position
- shadow map (depth buffer) is *not writeable during depth test*
- shadow mapping *does not discard fragments*

## Depth test setup for LDI generation:

- fragment must be **farther** away than fragment in previous depth layer → shadow map test
- fragment must be **the nearest** of all remaining fragments → regular depth test



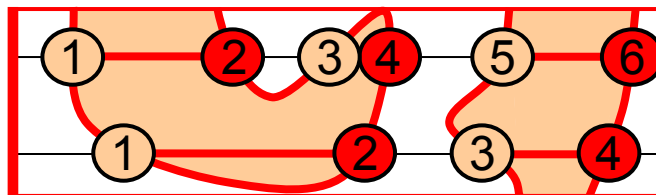
# Multipass LDI Generation





# *Result of LDI Generation*

- multipass LDI generation results in an ordered LDI representation of the Vol



Vol

1	2	3	4	5	6
1	2	3	4		

ordered LDI

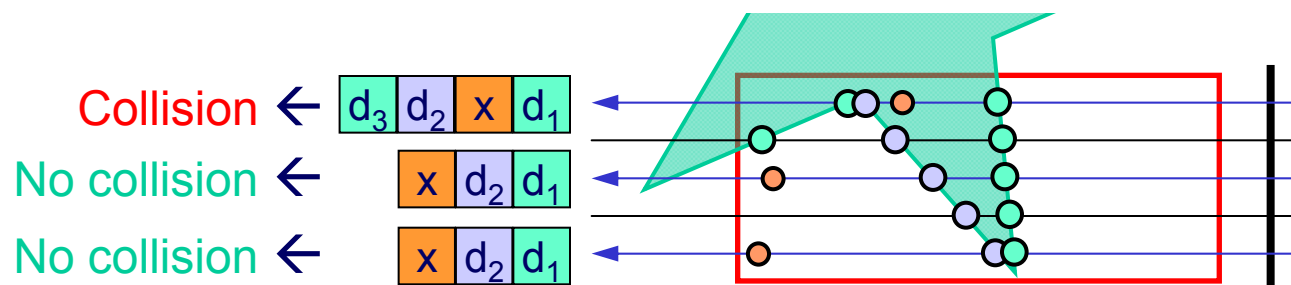
- requires one rendering pass per depth layer
- requires shadow mapping functionality



# *Collision Detection Test*

- test object primitives of one object against LDI of the other object (and vice versa)
- vertex-in-volume test

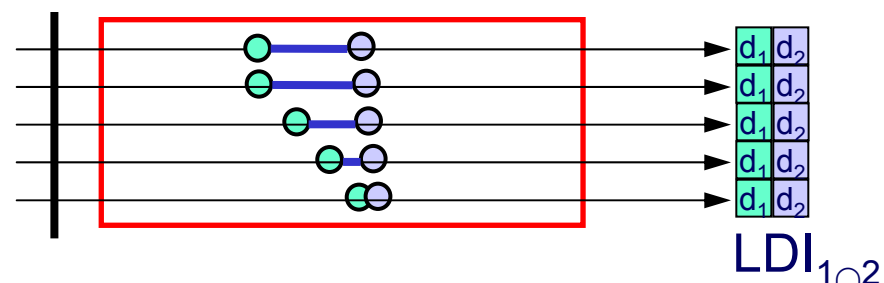
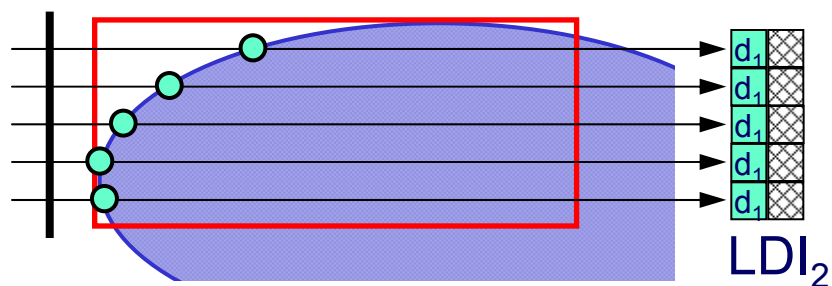
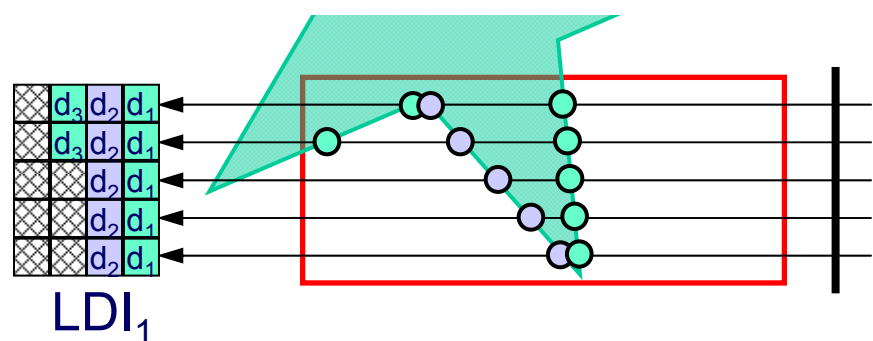
example:





# LDI Combination

- intersect both LDI to get the overlapping volume
- provides an explicit intersection volume
- other boolean operations (union, difference) are also possible  
→ constructive solid geometry (CSG)

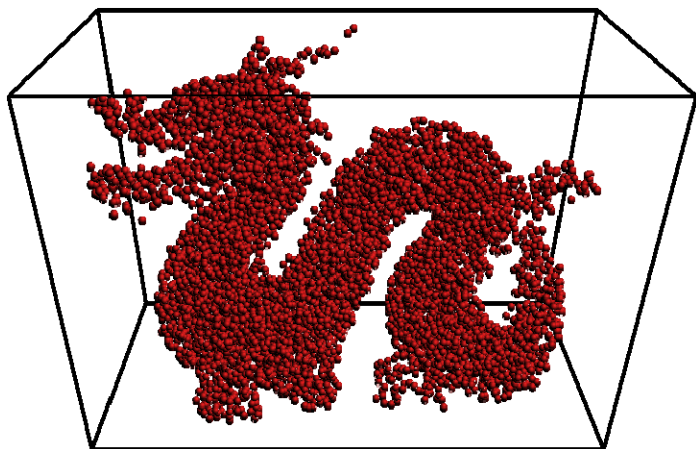


$$LDI_{1 \cap 2} = LDI_1 \cap LDI_2$$

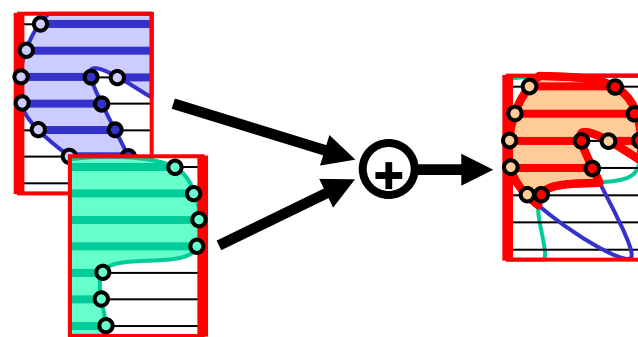
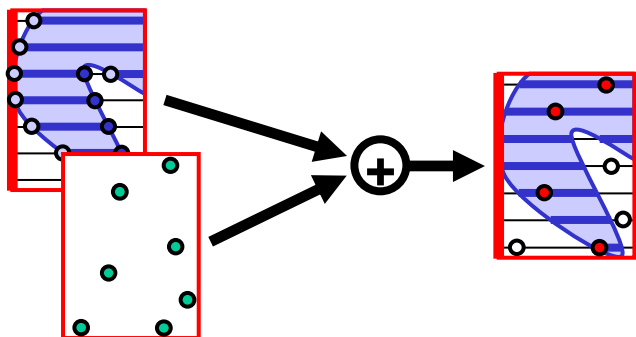
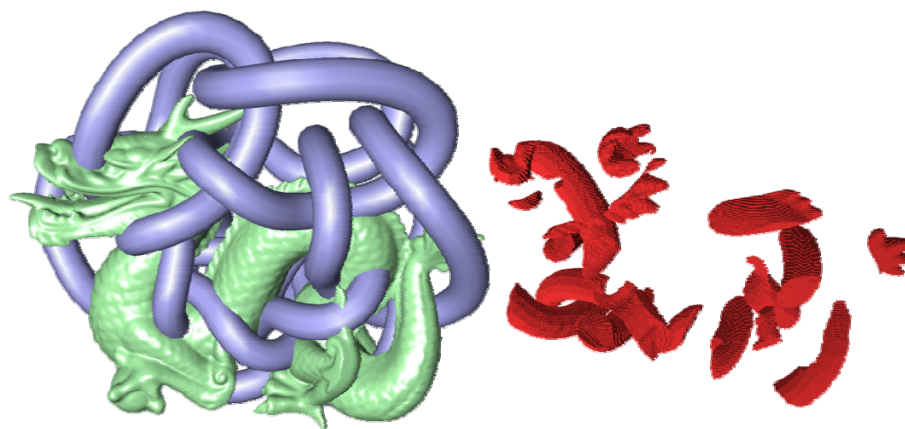


# *Collision queries*

## Vertex-in-volume test



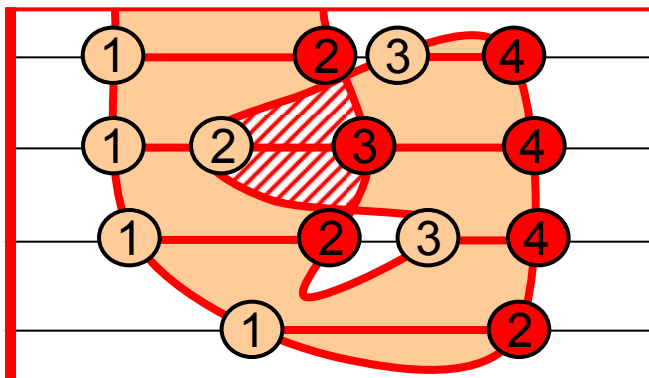
## Explicit intersection volume





# *Self-collision query*

- check for incorrect ordering of front and back faces
- if front and back faces do not alternate -> self collision



Vol

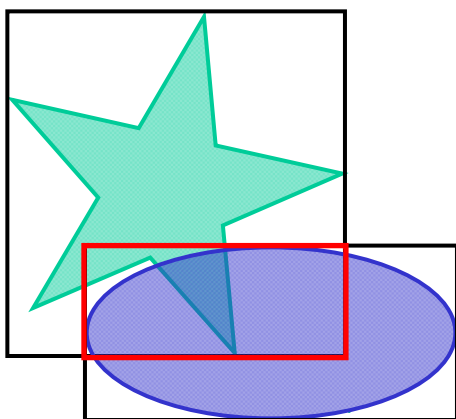
1	2	3	4
1	2	3	4
1	2	3	4
1	2		

LDI

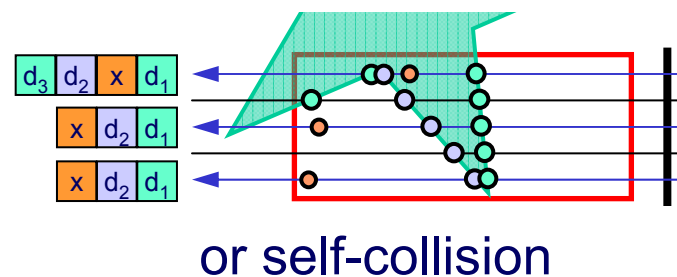
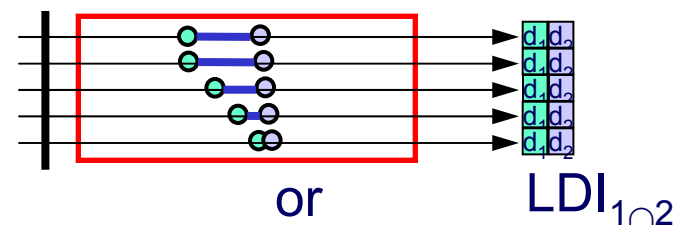


# Algorithm Summary

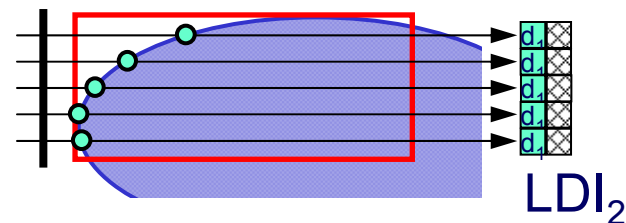
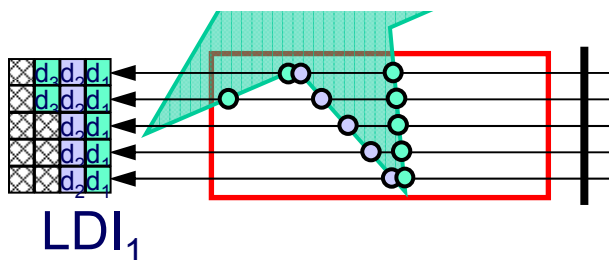
## (1) Volume of interest



## (3) Collision detection test



## (2) LDI generation



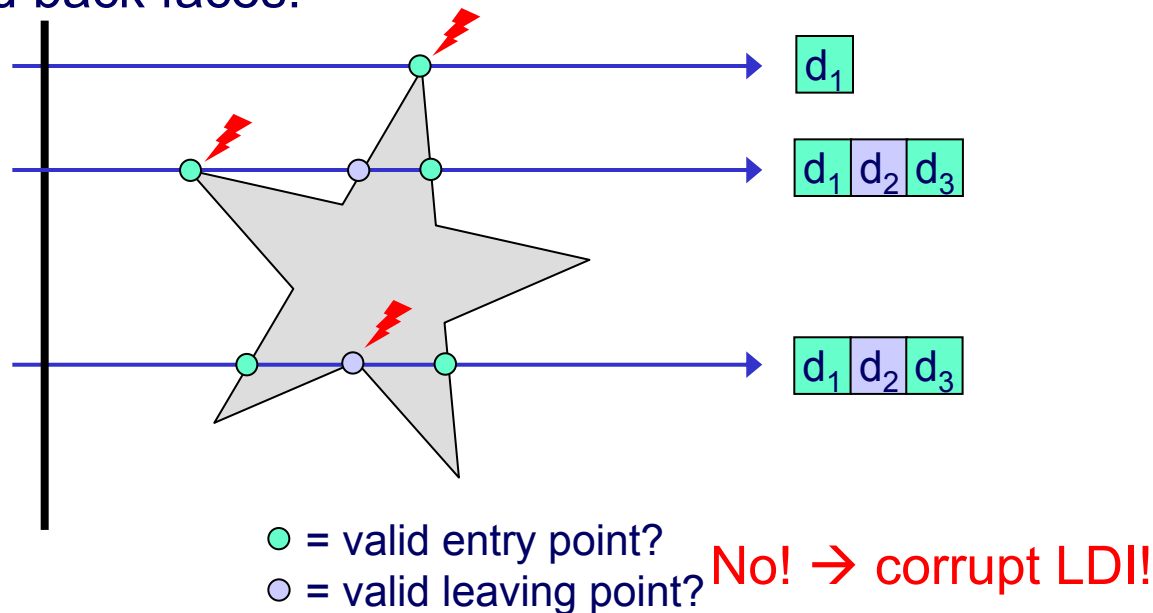




# Problems

- object can not be rendered to shadow map (see differences to depth buffer) → additional copy process necessary
- limited precision of depth buffer leads to singularities near edges between front and back faces:

example:

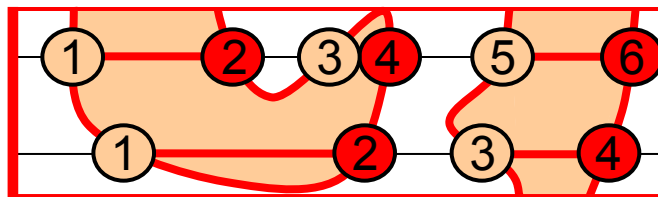


→ handle front and back faces in separate passes



# Unordered LDI Generation

- alternative method for LDI generation
- GPU generates unsorted LDI
  - fragments are rendered in the same order in each rendering pass
  - stencil buffer is used to get n-th value in the n-th pass
- CPU generates ordered LDI
  - depth complexity is known for each fragment  
(how many values are rendered per pixel)



Vol

5	3	2	1	4	6
4	1	3	2	2	2

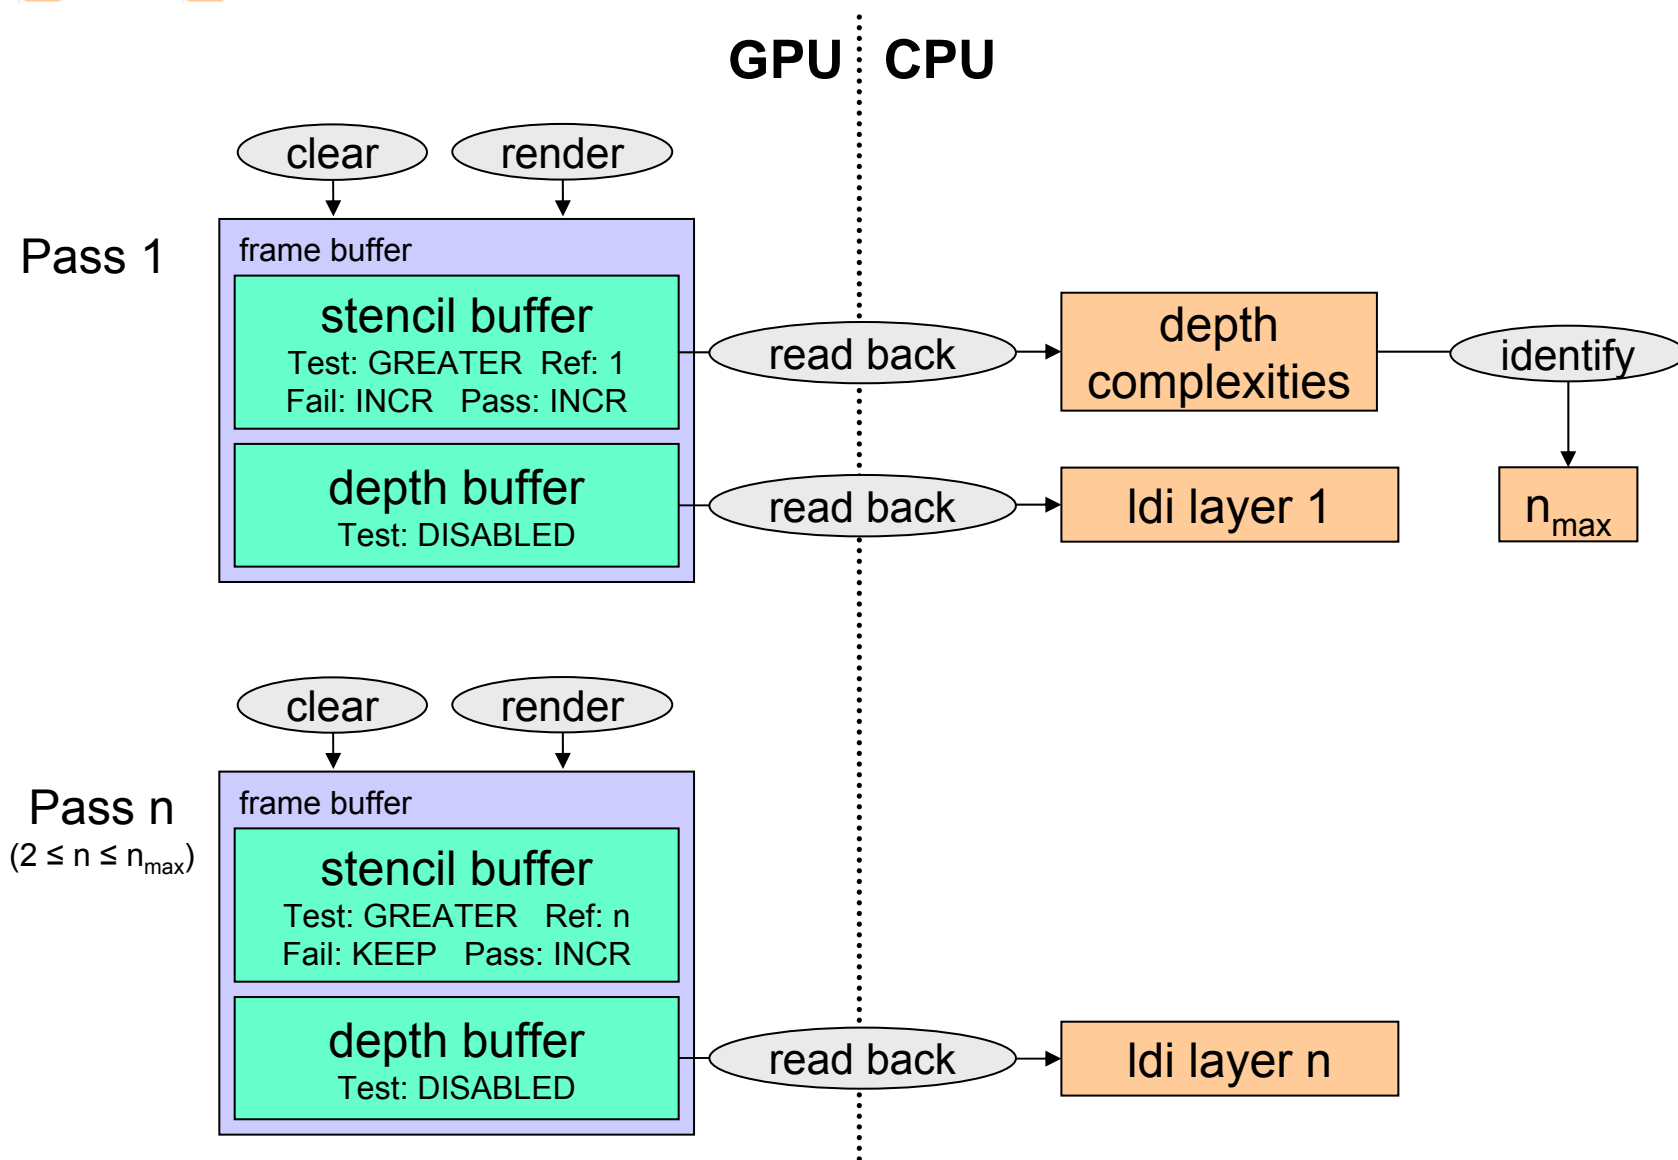
unsorted LDI (GPU)

1	2	3	4	5	6
1	2	3	4	2	2

sorted LDI (CPU)



# Unordered LDI Generation





# *Limitations*

- performance is dependent on:
  - depth complexity of objects in volume of interest
  - read back delay for simple objects
  - rendering speed for complex objects
- requires graphics hardware



# *Ordered LDI Generation on CPU*

## ***Motivation***

- buffer read-back from GPU  
can be performance bottleneck
- GPU requires multiple passes
- CPU can store fragments directly into LDI

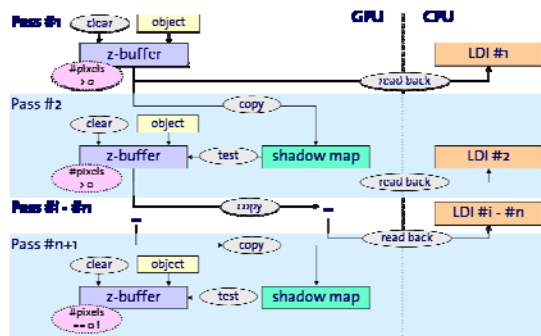
## ***Simplified software-renderer***

- rasterization of triangle meshes
- frustum culling
- face clipping
- orthogonal projection



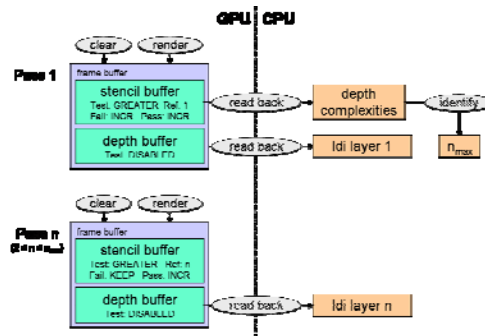
# LDI Generation - Summary

## Ordered LDI (GPU)



- $n+1$  passes
- complex setup
- two depth tests
- shadow map
- OpenGL extensions

## Unordered LDI (GPU)



- $n$  passes
- simple setup
- no depth test
- stencil buffer
- plain OpenGL 1.4

## Ordered LDI (CPU)

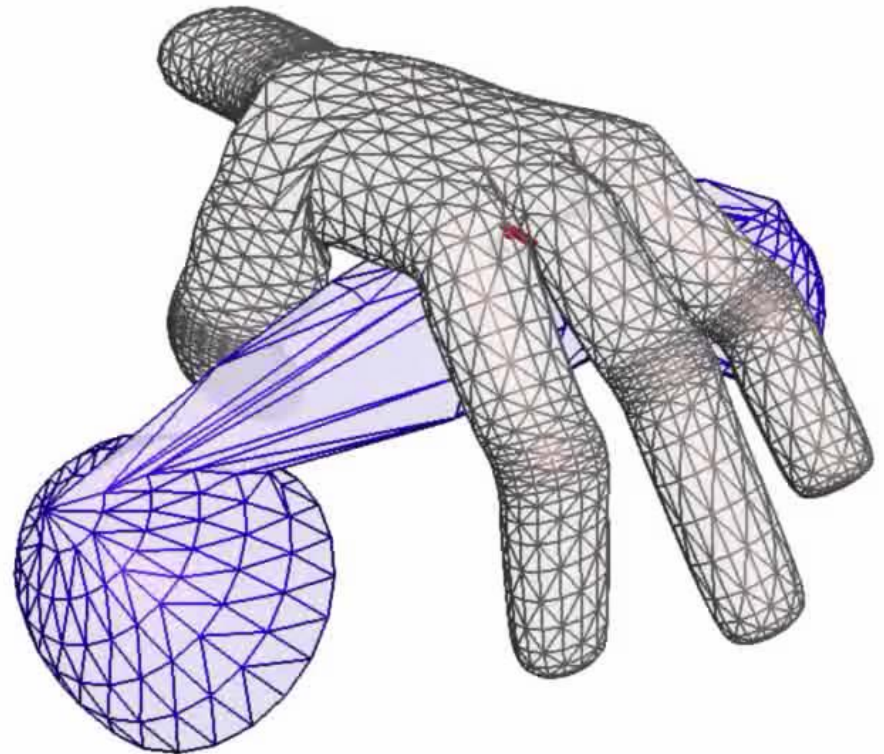
rasterize

- 1 pass
- simple setup
- no depth test



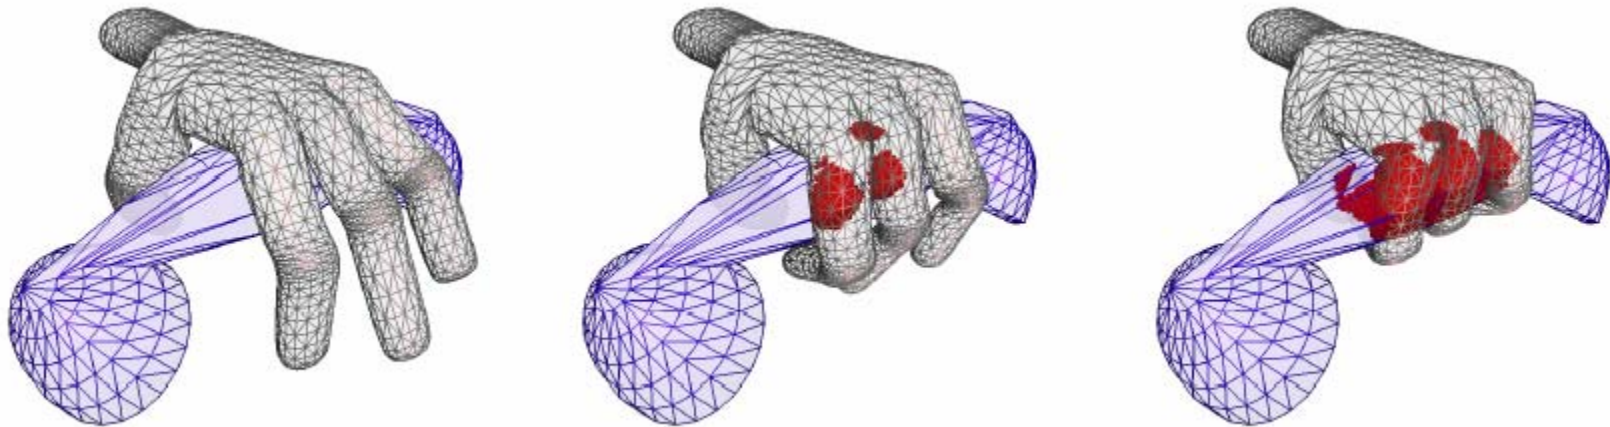
# *Performance - Intersection Volume*

- hand with 4800 faces
- phone with 900 faces
- two LDIs
- intersection volume for collision detection
- analysis of front / back face ordering for self-collision





# *Performance – Intersection Volume*



method	collision min / max	self collision min / max	overall min / max
ordered (GPU)	28 / 37	40 / 54	68 / 91
unordered (GPU, CPU)	9 / 12	12 / 18	21 / 30
software (CPU)	3 / 4	5 / 7	8 / 11

3 GHz Pentium 4, GeForce FX Ultra 5800

hand with 4800 faces  
phone with 900 faces  
measurements in ms





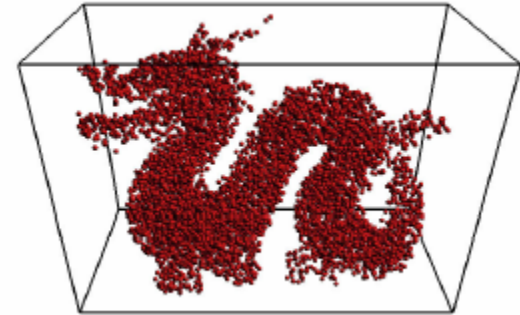
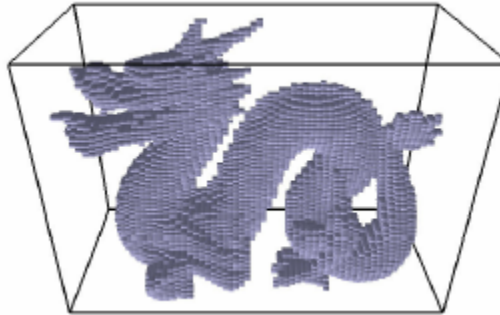
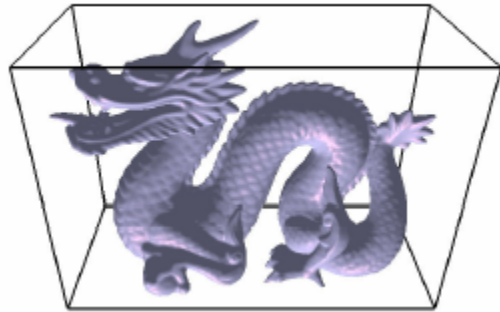
# *Performance – Vertex-in-Volume*

- santa with 10000 faces
- 20000 particles
- one LDI
- test vertices against  
inside regions of the LDI





# *Performance – Vertex-in-Volume*



method	520k faces 100k particles	150k faces 30k particles	50k faces 10k particles
ordered (GPU)	450	160	50
unordered (GPU, CPU)	225	75	25
software (CPU)	400	105	35

3 GHz Pentium 4, GeForce FX Ultra 5800

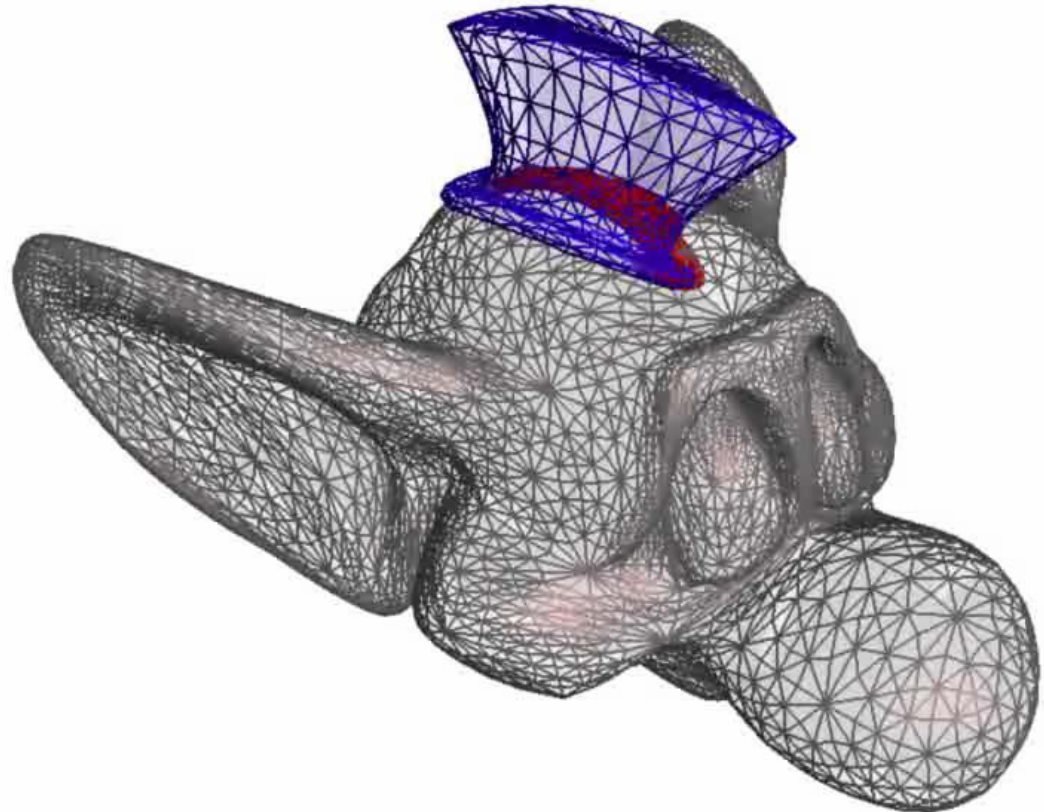
LDI resolution 64 x 64  
measurements in ms



# *Performance – LDI resolution*

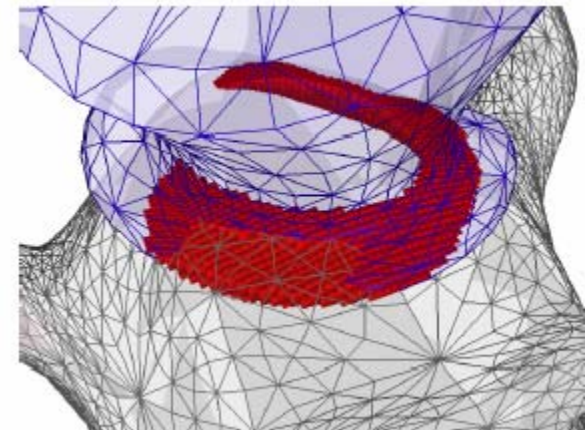
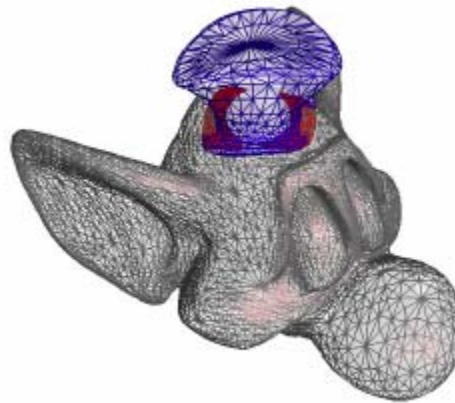
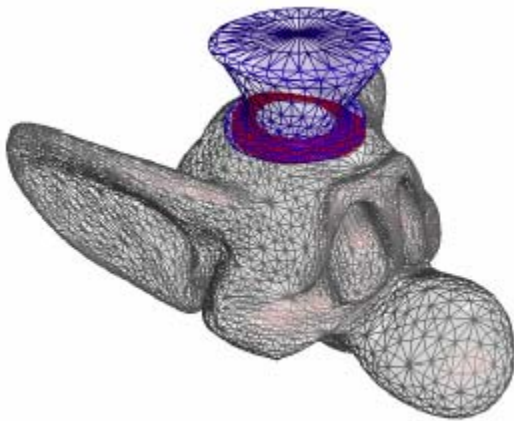
---

- mouse with 15000 faces
- hat with 1500 faces
- two LDIs
- intersection volume for collision detection





# *Performance – LDI resolution*



method	32 x32	64 x 64	128 x128
ordered (GPU)	24	26	51
unordered (GPU, CPU)	8	9	17
software (CPU)	2	3	6

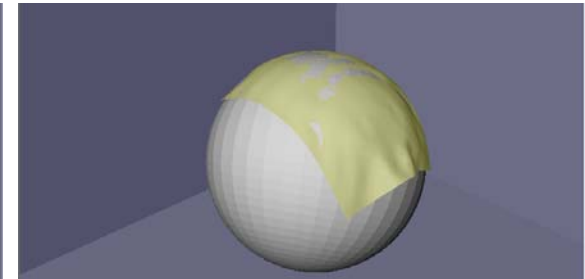
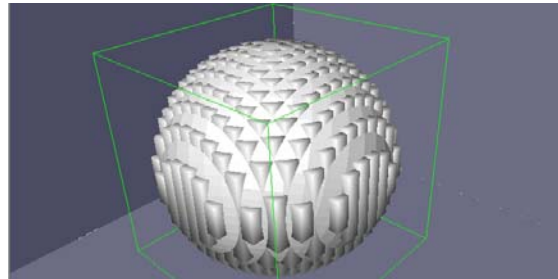
3 GHz Pentium 4, GeForce FX Ultra 5800

mouse with 15000 faces  
hat with 1500 faces  
measurements in ms

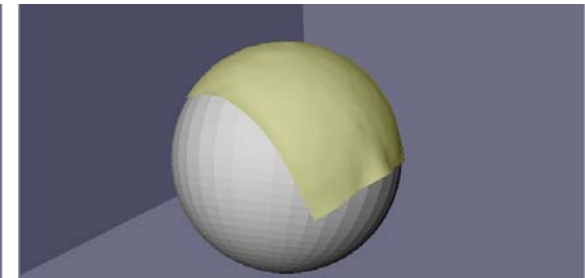
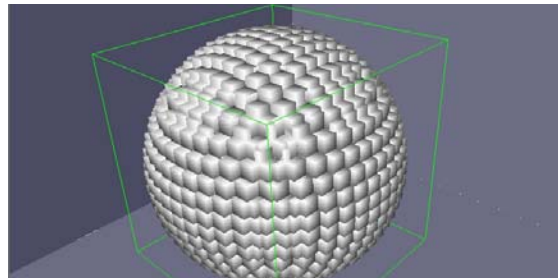


# *Applications – Cloth Modeling*

LDI



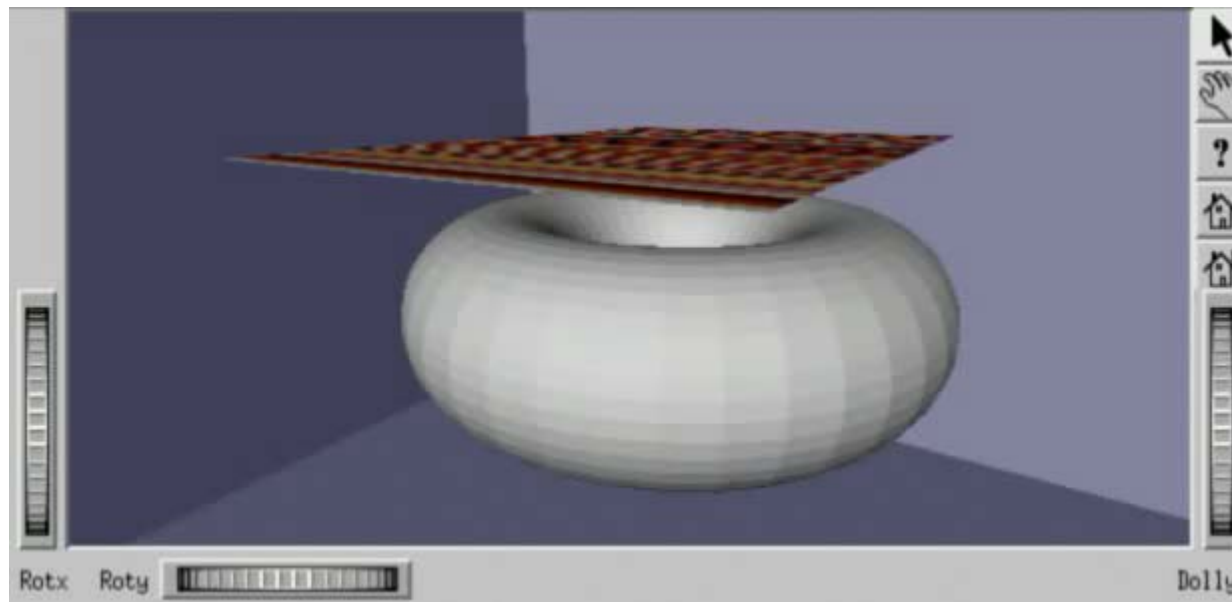
3 orthogonal  
dilated LDIs





# *Real-Time Cloth Simulation with Collision Handling*

real-time movie  
3GHz Pentium 4



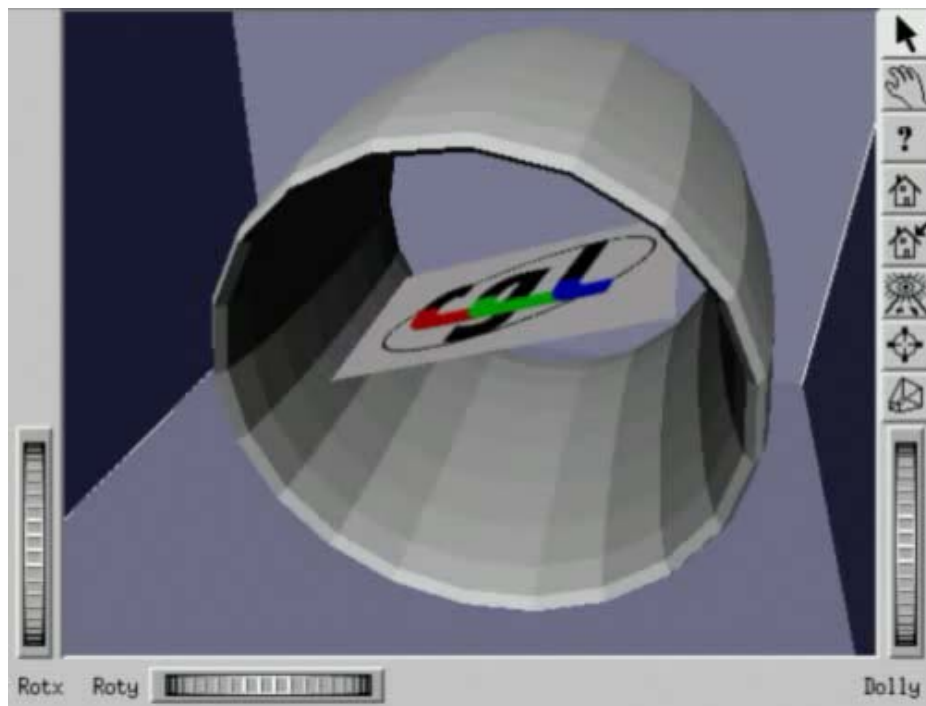
stable collision handling



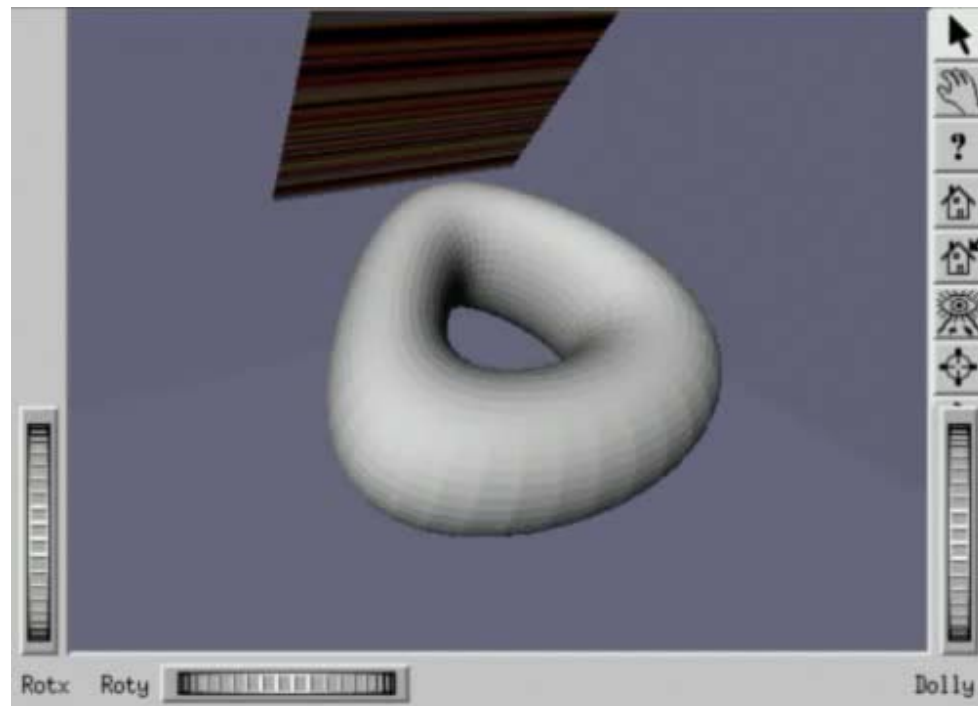


# *Real-Time Cloth Simulation with Collision Handling*

real-time movies  
3GHz Pentium 4



concave transforming object



concave deforming object



# *Summary*

- image-space technique
- detection of collisions and self-collisions
- handling of rigid and deformable closed meshes
- no pre-processing
- CPU: 5000 / 1000 faces at 100 Hz
- GPU: 520000 faces / 100000 particles at 4 Hz
- application to cloth simulation
- limitations
  - closed meshes
  - accuracy
  - collision information for collision response





# References

- M. Shinya, M. Forgue, “Interference Detection through rasterization,” *Journal of Visualization and Computer Animation*, vol. 2, pp. 132-134, 1991.
- K. Myszkowski, O. Okunev, T. Kunii, “Fast collision detection between complex solids using rasterizing graphics hardware,” *The Visual Computer*, vol. 11, no. 9, pp. 497-512, 1995.
- J. C. Lombardo, M.-P. Cani, F. Neyret, “Real-time Collision Detection for Virtual Surgery,” *Proc. of Comp. Anim.*, pp. 82-91, 1999.
- G. Baciú, S. K. Wong “Image-Based Techniques in a Hybrid Collision Detector,” *IEEE Trans on Visualization and Computer Graphics*, Jan 2002.
- D. Knott, D. Pai: “Cinder: Collision and interference detection in real-time using graphics hardware,” *Proc. Graphics Interface*, 2003.
- B. Heidelberger, M. Teschner, M. Gross, “Volumetric Collision Detection for Deformable Objects,” *Proc. VMV’03*, pp. 461-468, 2003.
- B. Heidelberger, M. Teschner, M. Gross, “Detection of Collisions and Self-collisions Using Image-space Techniques,” *Proc. WSCG’04*, pp. 145-152, 2004.